

## Introduction

This project encompasses an in-depth analysis and prediction model development for agricultural yield. The datasets which was utilized in this analysis include information regarding pesticides, temperature, rainfall, and historical yield data. It was gotten from Kaggle <https://www.kaggle.com/datasets/patelris/crop-yield-prediction-dataset>. The primary aim is to predict crop yield based on various factors such as temperature, pesticide usage, and environmental conditions. Although, we couldn't get data that speak to the Nigerian Ecosystem, we deem it necessary to use the available data to serve as a template for future work.

## Importing the necessary libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

## Reading the dataset

```
In [2]: df_pesticides = pd.read_csv("pesticides.csv")
df_yield_df = pd.read_csv("yield_df.csv")
df_yield = pd.read_csv("yield.csv")
df_temp = pd.read_csv("temp.csv")
df_rainfall = pd.read_csv("rainfall.csv")
```

## Data Overview

The datasets used in this study comprise:

Pesticides Data Yield Data Temperature Data Rainfall Data

```
In [3]: df_pesticides.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4349 entries, 0 to 4348
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Domain      4349 non-null   object
1   Area        4349 non-null   object
2   Element     4349 non-null   object
3   Item        4349 non-null   object
```

```

4   Year      4349 non-null   int64
5   Unit      4349 non-null   object
6   Value     4349 non-null   float64
dtypes: float64(1), int64(1), object(5)
memory usage: 238.0+ KB

```

In [4]:

```
df_yield.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 56717 entries, 0 to 56716
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Domain Code     56717 non-null  object
1   Domain          56717 non-null  object
2   Area Code       56717 non-null  int64
3   Area            56717 non-null  object
4   Element Code    56717 non-null  int64
5   Element         56717 non-null  object
6   Item Code       56717 non-null  int64
7   Item            56717 non-null  object
8   Year Code       56717 non-null  int64
9   Year            56717 non-null  int64
10  Unit            56717 non-null  object
11  Value           56717 non-null  int64
dtypes: int64(6), object(6)
memory usage: 5.2+ MB

```

In [5]:

```
df_temp.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 71311 entries, 0 to 71310
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   year        71311 non-null  int64
1   country     71311 non-null  object
2   avg_temp    68764 non-null  float64
dtypes: float64(1), int64(1), object(1)
memory usage: 1.6+ MB

```

In [6]:

```
df_rainfall.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6727 entries, 0 to 6726
Data columns (total 3 columns):
#   Column                                Non-Null Count  Dtype
---  ---
0   Area                                  6727 non-null   object
1   Year                                  6727 non-null   int64
2   average_rain_fall_mm_per_year       5953 non-null   object
dtypes: int64(1), object(2)
memory usage: 157.8+ KB

```

In [7]:

```
df_yield_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 28242 entries, 0 to 28241
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            28242 non-null  int64
1   Area                                  28242 non-null  object
2   Item                                  28242 non-null  object
3   Year                                  28242 non-null  int64
4   hg/ha_yield                           28242 non-null  int64
5   average_rain_fall_mm_per_year        28242 non-null  float64
6   pesticides_tonnes                     28242 non-null  float64
7   avg_temp                              28242 non-null  float64
dtypes: float64(3), int64(3), object(2)
memory usage: 1.7+ MB
```

**Inspecting the characteristics of the datasets, to know the available columns, unique values in certain columns, and potential columns for merging or joining datasets based on common their attributes.**

```
In [8]: print(df_yield.columns)
```

```
Index(['Domain Code', 'Domain', 'Area Code', 'Area', 'Element Code', 'Element',
      'Item Code', 'Item', 'Year Code', 'Year', 'Unit', 'Value'],
      dtype='object')
```

```
In [9]: print(df_yield['Area'].nunique())
print(df_yield['Year'].nunique())
```

```
212
56
```

```
In [10]: print(df_yield_df.columns)
```

```
Index(['Unnamed: 0', 'Area', 'Item', 'Year', 'hg/ha_yield',
      'average_rain_fall_mm_per_year', 'pesticides_tonnes', 'avg_temp'],
      dtype='object')
```

```
In [11]: print(df_yield_df['Area'].nunique())
print(df_yield_df['Year'].nunique())
df_yield_df.drop(['Unnamed: 0'], axis = 1 ,inplace = True)
```

```
101
23
```

```
In [12]: print(df_rainfall.columns)
```

```
Index(['Area', 'Year', 'average_rain_fall_mm_per_year'], dtype='object')
```

```
In [13]: # Renaming the 'Area' column to remove the space before it for uniformity
df_rainfall.rename(columns = {'Area': 'Area'}, inplace = True) #

print(df_rainfall['Area'].nunique())
print(df_rainfall['Year'].nunique())
```

217  
31

In [14]:

```
print(df_temp.columns)
```

```
Index(['year', 'country', 'avg_temp'], dtype='object')
```

In [15]:

```
# Renaming the ' year' column to 'Year' for uniformity
df_temp.rename(columns = {'year':'Year','country':'Area'},inplace = True)

print(df_temp['Area'].nunique())
print(df_temp['Year'].nunique())
print(df_temp.columns)
```

137

271

```
Index(['Year', 'Area', 'avg_temp'], dtype='object')
```

**From the Analysis Above, it is clear that the Year and and Country(Area) columns are common to all , this will be used for merging the datasets on these columns.**

In [16]:

```
df_temprain = pd.merge(df_rainfall,df_temp,on = ['Year','Area'])
df_temprainpest = pd.merge(df_temprain,df_pesticides , on = ['Year','Area'])
```

In [17]:

```
#Find common columns between df_trp and df_yield by iterating through their columns.
trplist =list(df_temprainpest.columns)
dfyield = list(df_yield.columns)
# Create a list (com1) containing column names present in both DataFrames.
com1 = [i for i in trplist if i in dfyield]
```

In [18]:

```
print(df_yield.columns)
print(df_yield_df.columns)
```

```
ly = [i for i in list(df_yield.columns) if i in (df_yield_df.columns)]
```

```
Index(['Domain Code', 'Domain', 'Area Code', 'Area', 'Element Code', 'Element',
      'Item Code', 'Item', 'Year Code', 'Year', 'Unit', 'Value'],
      dtype='object')
Index(['Area', 'Item', 'Year', 'hg/ha_yield', 'average_rain_fall_mm_per_year',
      'pesticides_tonnes', 'avg_temp'],
      dtype='object')
```

In [19]:

```
yield_df = pd.merge(df_yield_df,df_yield, on = ['Year','Area','Item'])
```

In [20]:

```
print(yield_df.shape)

print(yield_df.columns)
print(df_temprainpest.columns)
```

```
(28242, 16)
Index(['Area', 'Item', 'Year', 'hg/ha_yield', 'average_rain_fall_mm_per_year',
      'pesticides_tonnes', 'avg_temp', 'Domain Code', 'Domain', 'Area Code',
      'Element Code', 'Element', 'Item Code', 'Year Code', 'Unit', 'Value'],
      dtype='object')
Index(['Area', 'Year', 'average_rain_fall_mm_per_year', 'avg_temp', 'Domain',
      'Element', 'Item', 'Unit', 'Value'],
      dtype='object')
```

```
In [21]: #Assuming your first dataframe is df1 and second dataframe is df2
years_df1 = set(yield_df['Year'].unique())
years_df2 = set(df_temprainpest['Year'].unique())

# Find the years in df2 that are not in df1
years_only_in_df2 = years_df2 - years_df1
```

```
In [22]: # Print the result
print("Years present in df1 but not in df2:", years_only_in_df2)
```

Years present in df1 but not in df2: set()

This result shows that are no mismatch or missing years in terms of the 'Year' column when comparing the two dataframes.

```
In [23]: area_df1 = set(yield_df['Year'].unique())
area_df2 = set(df_temprainpest['Year'].unique())
```

```
In [24]: # Doing same for Area, checking if there are Areas in df2 that are not in df1
area_only_in_df1 = area_df2 - area_df1

# Print the result
print("Area present in df1 but not in df2:", area_only_in_df1)
```

Area present in df1 but not in df2: set()

## EDA

```
In [25]: yield_df.head()
```

```
Out[25]:
```

	Area	Item	Year	hg/ha_yield	average_rain_fall_mm_per_year	pesticides_tonnes	avg_temp	Do
0	Albania	Maize	1990	36613	1485.0	121.0	16.37	
1	Albania	Potatoes	1990	66667	1485.0	121.0	16.37	
2	Albania	Rice, paddy	1990	23333	1485.0	121.0	16.37	
3	Albania	Sorghum	1990	12500	1485.0	121.0	16.37	
4	Albania	Soybeans	1990	7000	1485.0	121.0	16.37	

```
In [26]: print(yield_df['Item Code'].nunique())
print(yield_df['Item'].nunique())
## It has one unique value for each crop

10
10
```

```
In [27]: print(yield_df['Domain'].nunique())
print(yield_df['Domain Code'].nunique())

1
1
```

```
In [28]: print(yield_df['Area'].nunique())
print(yield_df['Area Code'].nunique())

101
101
```

```
In [29]: print(yield_df['Element Code'].nunique())
print(yield_df['Element'].nunique())
print(yield_df['Unit'].nunique())

1
1
1
```

Based on these information, it is clear that the 'Year Code' and 'Domain Code' columns are redundant and unnecessary. The 'Area Code' column does not contribute meaningfully to drawing conclusions. Since the 'Domain' column has only one unique value, it is deemed irrelevant and is dropped. Additionally, 'Area Code' is mentioned twice, implying redundancy. The columns 'hg/ha\_yield', 'Value', and 'Unit' are interrelated, and as a result, the last two columns are being dropped.

```
In [30]: Yield_dfcopy = yield_df.copy()
```

```
In [31]: Yield_dfcopy.drop(['Area Code', 'Year Code', 'Domain', 'Domain Code',
                          'Area Code', 'Item Code', 'Element', 'Element Code', 'Unit', 'Value'],
                          axis = 1, inplace = True)
```

```
In [32]: Yield_dfcopy.head()
```

```
Out[32]:
```

	Area	Item	Year	hg/ha_yield	average_rain_fall_mm_per_year	pesticides_tonnes	avg_temp
0	Albania	Maize	1990	36613	1485.0	121.0	16.37
1	Albania	Potatoes	1990	66667	1485.0	121.0	16.37
2	Albania	Rice, paddy	1990	23333	1485.0	121.0	16.37
3	Albania	Sorghum	1990	12500	1485.0	121.0	16.37
4	Albania	Soybeans	1990	7000	1485.0	121.0	16.37

```
In [33]: Yield_dfcopy[Yield_dfcopy['Area'] == 'Ghana'].head()
```

```
Out[33]:
```

	Area	Item	Year	hg/ha_yield	average_rain_fall_mm_per_year	pesticides_tonnes	avg_temp
<b>9161</b>	Ghana	Cassava	1990	84170	1187.0	65.8	26.73
<b>9162</b>	Ghana	Maize	1990	11889	1187.0	65.8	26.73
<b>9163</b>	Ghana	Plantains and others	1990	61890	1187.0	65.8	26.73
<b>9164</b>	Ghana	Rice, paddy	1990	16510	1187.0	65.8	26.73
<b>9165</b>	Ghana	Sorghum	1990	6310	1187.0	65.8	26.73

The dataset should include the average temperature in Celsius, calculated based on temperature values.

## Feature Engineering: Splitting Items with Multiple Crops

The 'Item' column in Yield\_final\_data contains entries where multiple crops are listed using commas. To address this, the column is split using the comma as a delimiter, and the resulting lists are expanded into separate rows using the explode function. The index is then reset for clarity.

```
In [34]: Yield_dfcopy['Item'] = Yield_dfcopy['Item'].str.split(',')
Yield_dfcopy = Yield_dfcopy.explode('Item').reset_index(drop=True)
```

```
In [35]: Yield_dfcopy.shape
```

```
Out[35]: (31630, 7)
```

```
In [36]: Yield_dfcopy.loc[Yield_dfcopy['Area'].str.len().sort_values().index].head(12)
```

```
Out[36]:
```

	Area	Item	Year	hg/ha_yield	average_rain_fall_mm_per_year	pesticides_tonnes	avg_temp
<b>17260</b>	Iraq	Potatoes	1996	177576	216.0	616.00	23.96
<b>17377</b>	Iraq	paddy	2005	28879	216.0	873.01	19.96
<b>17376</b>	Iraq	Rice	2005	28879	216.0	873.01	19.96
<b>17375</b>	Iraq	paddy	2005	28879	216.0	873.01	23.58
<b>17374</b>	Iraq	Rice	2005	28879	216.0	873.01	23.58
<b>17373</b>	Iraq	Potatoes	2005	158431	216.0	873.01	19.96
<b>17372</b>	Iraq	Potatoes	2005	158431	216.0	873.01	23.58
<b>17371</b>	Iraq	Maize	2005	23079	216.0	873.01	19.96

	Area	Item	Year	hg/ha_yield	average_rain_fall_mm_per_year	pesticides_tonnes	avg_temp
17370	Iraq	Maize	2005	23079	216.0	873.01	23.58
17369	Iraq	Wheat	2004	11898	216.0	828.62	20.00
17368	Iraq	Wheat	2004	11898	216.0	828.62	23.72
17367	Iraq	Soybeans	2004	10000	216.0	828.62	20.00

In [37]:

Yield\_dfcopy['Item'].value\_counts()

Out[37]:

Potatoes4276  
Maize4121  
Wheat3857  
Rice3388  
paddy3388  
Soybeans3223  
Sorghum3039  
Sweet potatoes2890  
Cassava2045  
Yams847  
Plantains and others556  
Name: Item, dtype: int64

In [38]:

Yield\_dfcopy.info()

<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 31630 entries, 0 to 31629  
Data columns (total 7 columns):  
# Column Non-Null Count Dtype  
--- -  
0 Area 31630 non-null object  
1 Item 31630 non-null object  
2 Year 31630 non-null int64  
3 hg/ha\_yield 31630 non-null int64  
4 average\_rain\_fall\_mm\_per\_year 31630 non-null float64  
5 pesticides\_tonnes 31630 non-null float64  
6 avg\_temp 31630 non-null float64  
dtypes: float64(3), int64(2), object(2)  
memory usage: 1.7+ MB

In [39]:

Yield\_dfcopy.describe()

Out[39]:

	Year	hg/ha_yield	average_rain_fall_mm_per_year	pesticides_tonnes	avg_temp
count	31630.000000	31630.000000	31630.000000	31630.000000	31630.000000
mean	2001.536895	73162.659437	1153.085267	37062.481868	20.620043
std	7.053614	81302.121833	712.805826	59771.245918	6.255367
min	1990.000000	50.000000	51.000000	0.040000	1.300000
25%	1995.000000	20918.000000	593.000000	1714.390000	16.780000
50%	2001.000000	37607.000000	1083.000000	17529.440000	21.550000
75%	2008.000000	94651.000000	1668.000000	48715.510000	26.030000



	Year	hg/ha_yield	average_rain_fall_mm_per_year	pesticides_tonnes	avg_temp
<b>max</b>	2013.000000	501412.000000	3240.000000	367778.000000	30.650000

```
In [40]: Yield_dfcopy.groupby(['Area'], sort = True)[['hg/ha_yield']].sum().nlargest(10, 'hg/ha_y
```

```
Out[40]:
```

	hg/ha_yield
<b>Area</b>	
<b>India</b>	342749968
<b>Brazil</b>	175935463
<b>Mexico</b>	139270288
<b>Japan</b>	133321962
<b>Australia</b>	121009524
<b>Pakistan</b>	80118288
<b>Indonesia</b>	75507330
<b>Turkey</b>	59531760
<b>United Kingdom</b>	55419990
<b>Spain</b>	51622833

```
In [41]: Yield_dfcopy.groupby(['Item', 'Area'], sort=True)[['hg/ha_yield']].sum().nlargest(20)
```

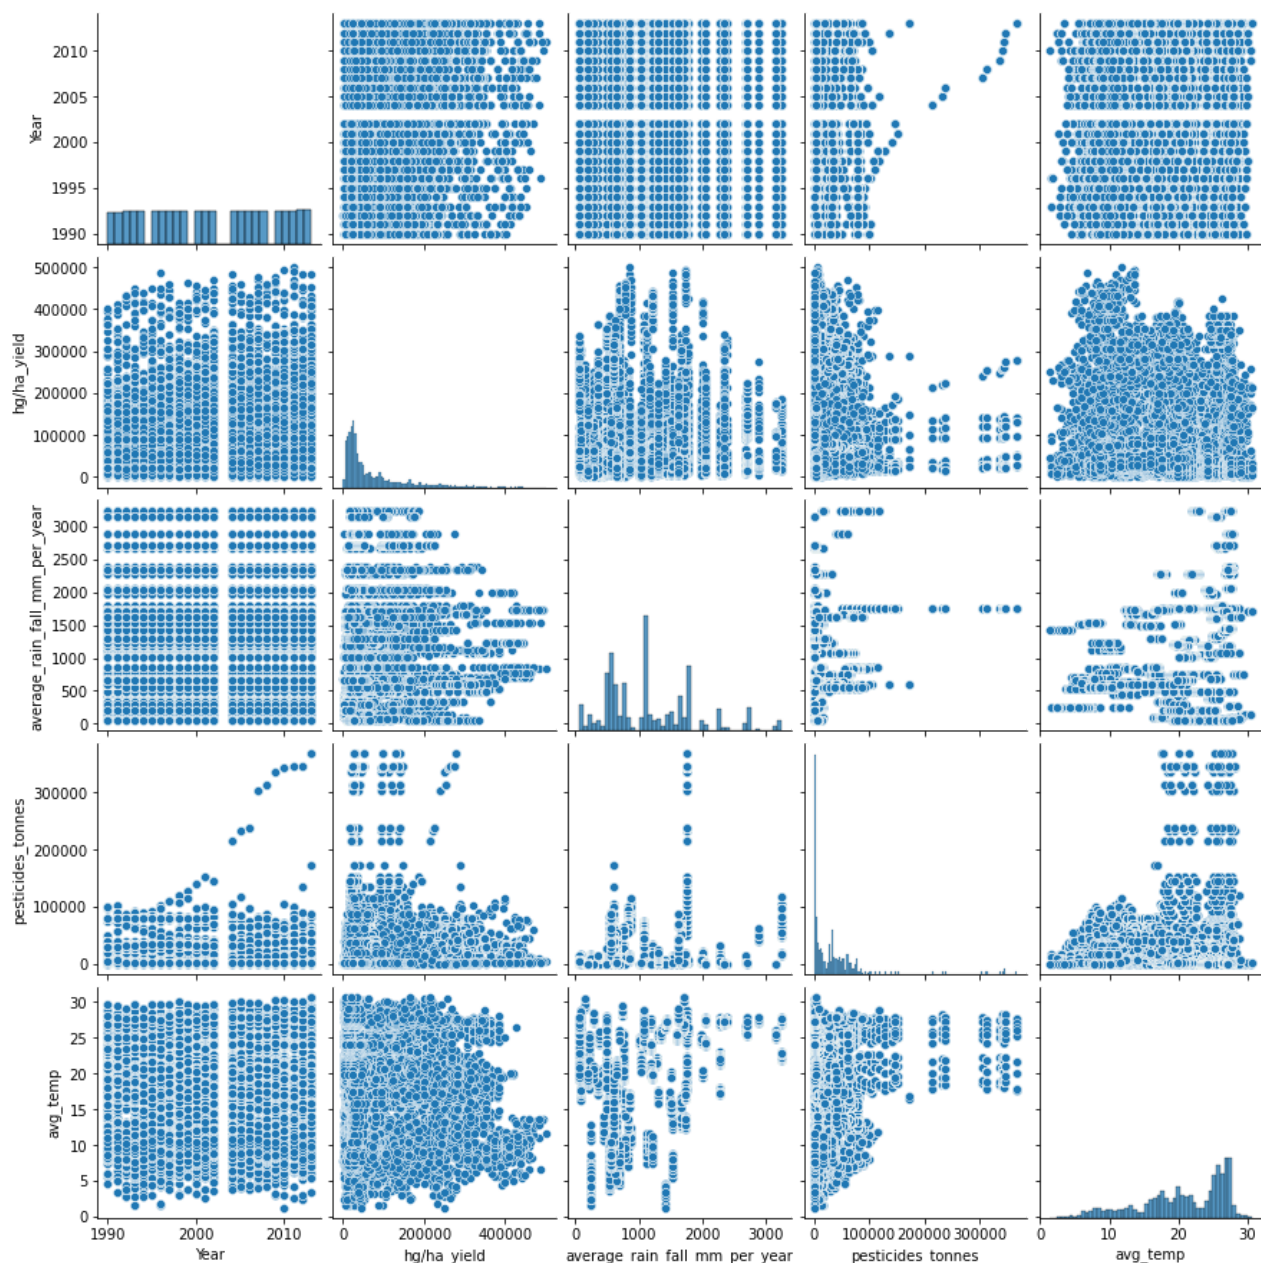
```
Out[41]:
```

Item	Area	
Cassava	India	142810624
Potatoes	India	92122514
	Brazil	49602168
	United Kingdom	46705145
	Australia	45670386
Sweet potatoes	India	44439538
Potatoes	Japan	42918726
	Mexico	42053880
Sweet potatoes	Mexico	35808592
	Australia	35550294
Cassava	Brazil	33671231
Potatoes	Pakistan	32969754
Sweet potatoes	Japan	32794236
Potatoes	Turkey	30530955
Yams	Japan	29165394
Sweet potatoes	Brazil	28266502
Potatoes	South Africa	27341980
	Germany	26672181
Yams	Brazil	23472053
Sweet potatoes	Pakistan	21687615

Name: hg/ha\_yield, dtype: int64

```
In [42]: ## Checking the pairplot between the columns  
sns.pairplot(Yield_dfcopy)
```

Out[42]: <seaborn.axisgrid.PairGrid at 0x22af69b8550>



```
In [43]: # Checking the data for outliers

plt.figure(figsize = (10,10))

plt.subplot(3,2,1)
sns.boxplot(data= Yield_dfcopy['Year'])
plt.title('Year')

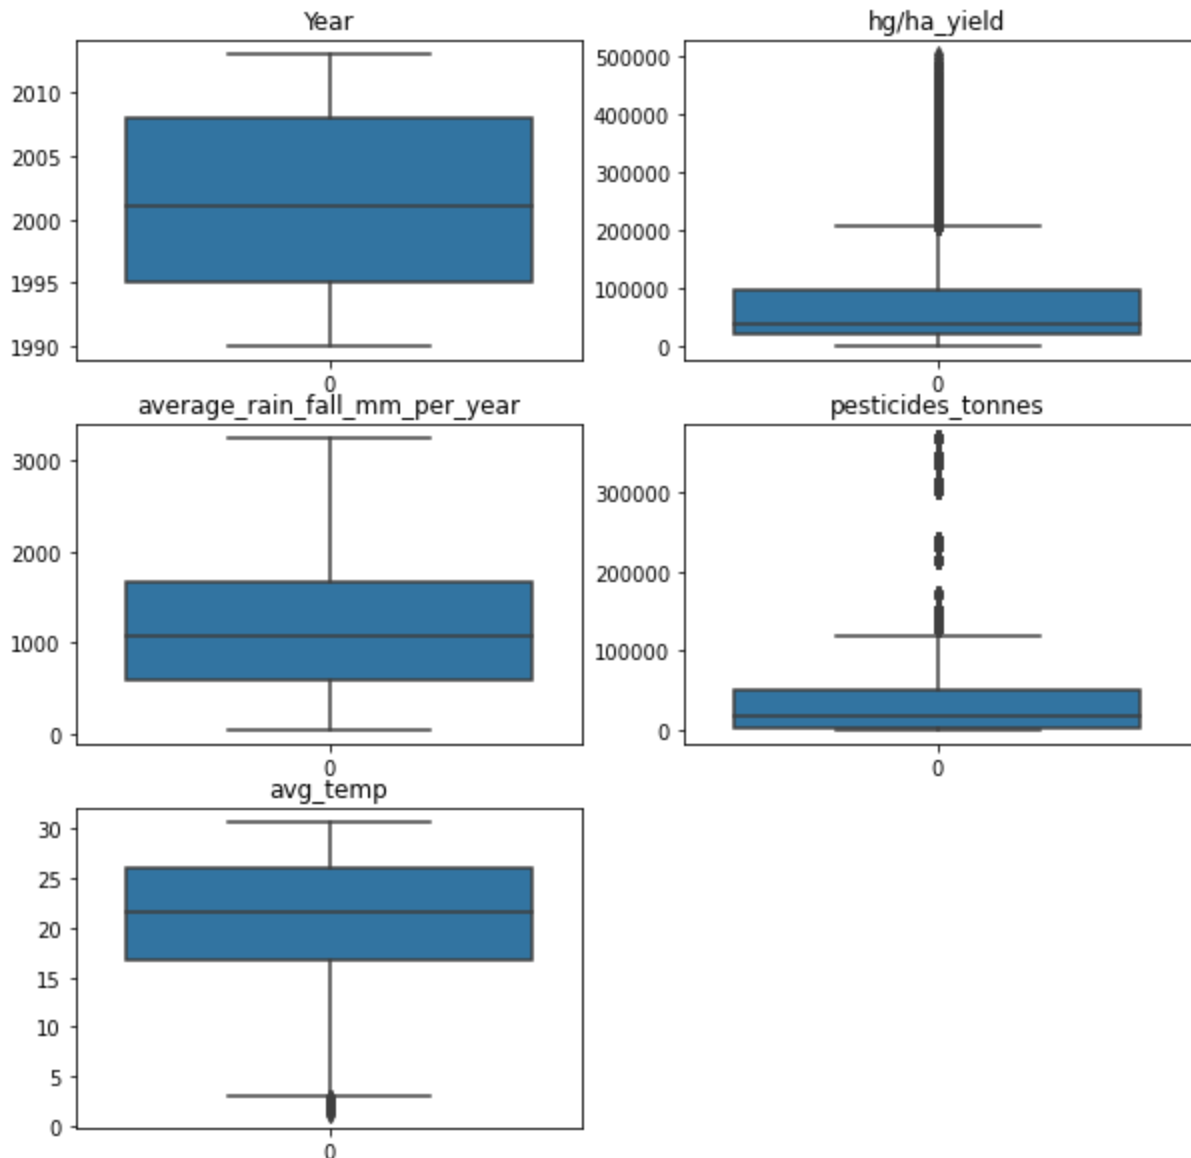
plt.subplot(3,2,2)
sns.boxplot(data= Yield_dfcopy['hg/ha_yield'])
plt.title('hg/ha_yield')

plt.subplot(3,2,3)
sns.boxplot(data= Yield_dfcopy['average_rain_fall_mm_per_year'])
plt.title('average_rain_fall_mm_per_year')
```

```
plt.subplot(3,2,4)
sns.boxplot(data= Yield_dfcopy['pesticides_tonnes'])
plt.title('pesticides_tonnes')

plt.subplot(3,2,5)
sns.boxplot(data= Yield_dfcopy['avg_temp'])
plt.title('avg_temp')

plt.show()
```



It has been observed that there are outliers in the 'avg\_temp' and 'pesticides\_tonnes' columns

In [ ]:


In [44]:

```
num_cor = Yield_dfcopy.select_dtypes(['int64', 'float64']).corr()
```

In [45]:

```
sns.heatmap(num_cor, cmap = 'YlGnBu', annot = True)
plt.title('Heatmap')
```


Out[45]: Text(0.5, 1.0, 'Heatmap')



The plot visually indicates that there is no significant correlation observed among the columns. The absence of strong correlations suggests that the variables may not have a linear relationship or exhibit a direct influence on each other. Further statistical analysis may be conducted to explore any underlying patterns or relationships.


In [46]: `#Checking the data distribution in the  
sns.histplot(Yield_dfcopy, x = 'Year' , bins = range(1985, 2020, 5))`

Out[46]: <AxesSubplot:xlabel='Year', ylabel='Count'>



In [47]: `### Checking the pesticide usage data  
  
## Checking the data distribution in the yield column  
sns.histplot(Yield_dfcopy, x = 'pesticides_tonnes' )`

Out[47]: <AxesSubplot:xlabel='pesticides\_tonnes', ylabel='Count'>




There are outlier points in the datasets, to maintain data integrity while addressing the impact of extreme values on the analysis, mitigating potential high variance by removing only the upper 10% of data points is needed.

```
In [48]: Yield_dfcopy = Yield_dfcopy[Yield_dfcopy['pesticides_tonnes']
        <= Yield_dfcopy['pesticides_tonnes'].quantile(0.90)]
```


```
In [49]: # Checking the data distribution in the yield column(if there have been any changes in
sns.histplot(Yield_dfcopy, x = 'hg/ha_yield' )
```

Out[49]: <AxesSubplot:xlabel='hg/ha\_yield', ylabel='Count'>



```
In [50]: sns.histplot(Yield_dfcopy, x = 'pesticides_tonnes' )
```

Out[50]: <AxesSubplot:xlabel='pesticides\_tonnes', ylabel='Count'>



```
In [51]: Yield_dfcopy= pd.get_dummies(Yield_dfcopy,columns = ['Item','Area'], drop_first = True
```

```
In [ ]:
```

```
In [ ]:
```

## Splitting the data

```
In [52]: split_ratio = 0.25
split_index = int(len(Yield_dfcopy) * split_ratio)

# Get the 'Year' value at the split point
split_year = Yield_dfcopy['Year'].iloc[split_index]

print(f"Split Year: {split_year}")
```

Split Year: 2006

```
In [53]: df_train = Yield_dfcopy[Yield_dfcopy['Year'] <= 2008]
df_test = Yield_dfcopy[Yield_dfcopy['Year'] > 2008]
```

```
In [54]: Yield_dfcopy.shape
```

```
Out[54]: (28959, 115)
```

```
In [55]: df_train = df_train.drop('Year',axis = 1)
df_test = df_test.drop('Year',axis = 1)

X_train = df_train.drop('hg/ha_yield',axis = 1 )
y_train = df_train['hg/ha_yield']
X_test = df_test.drop('hg/ha_yield', axis = 1)
y_test = df_test['hg/ha_yield']

print(X_train.shape)
```

```
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)
```

```
(22632, 113)
(22632,)
(6327, 113)
(6327,)
```

```
In [56]: y_train.values.reshape(-1,1)
         y_test.values.reshape(-1,1)
```

```
Out[56]: array([[ 55693],
                [219780],
                [ 16667],
                ...,
                [ 13142],
                [ 22222],
                [ 22888]], dtype=int64)
```

## Scaling and Model training

```
In [57]: scaler = MinMaxScaler()
         X_train = scaler.fit_transform(X_train)
         X_test = scaler.transform(X_test)
```

## Training: Training of model using Linear Regression

```
In [59]: from sklearn.linear_model import LinearRegression
         LR = LinearRegression()
         LR.fit(X_train,y_train)
         y_pred = LR.predict(X_test)

         Model_perf = pd.DataFrame(columns=['Model_Name', 'MSE', 'R2_Score'])

         LR_mse = mean_squared_error(y_test,y_pred)
         LR_R2 = r2_score(y_test,y_pred)

         new_row = {'Model_Name':'Linear Regression', 'MSE':LR_mse , 'R2_Score': LR_R2}
         Model_perf = Model_perf.append(new_row,ignore_index = True)
```

```
In [60]: from sklearn.preprocessing import PolynomialFeatures

         # Create polynomial features
         poly_features = PolynomialFeatures(degree=1)
         X_poly = poly_features.fit_transform(X_train)

         # Train the polynomial regression model
         poly_regression = LinearRegression()
         poly_regression.fit(X_poly, y_train)

         # Predict using the trained model
```

```
X_test_poly = poly_features.transform(X_test)
y_pred = poly_regression.predict(X_test_poly)
print("Predicted values:", y_pred)
```

```
Predicted values: [ 22948.8125 174228.8125   2213.5625 ... -19339.5625  85013.875
-12284.5625]
```

```
In [61]: PR_mse = mean_squared_error(y_test,y_pred)
PR_R2 = r2_score(y_test,y_pred)

new_row = {'Model_Name':'Polynomial Regression(degree 1)', 'MSE':PR_mse , 'R2_Score': PR
Model_perf = Model_perf.append(new_row,ignore_index = True)
```

```
In [ ]:
```

```
In [ ]:
```

## Model Performance

```
In [62]: Model_perf
```

```
Out[62]:
```

	Model_Name	MSE	R2_Score
0	Linear Regression	7.542342e+28	-8.749149e+18
1	Polynomial Regression(degree 1)	3.652341e+25	-4.236731e+15

## Conclusion

This project aims to harness the power of data science to gain insights into agricultural trends, optimize crop yield predictions, and contribute to informed decision-making in the agricultural sector. Through careful analysis and modeling, we seek to provide valuable tools for farmers, policymakers, and stakeholders to enhance sustainable and efficient agricultural practices. Although, we couldn't get data that speak to the Nigerian Ecosystem, we deem it necessary to use the available data to serve as a template for future work.

## Recommendation

**Acknowledging the limitation in obtaining data specific to the Nigerian ecosystem, leveraging the available dataset becomes pivotal as a template for future endeavors within the Nigerian agricultural landscape. Despite the absence of direct data, this project serves as a foundational framework, offering insights into methodologies, potential analyses, and model development that can be adapted and fine-tuned when specific Nigerian agricultural data becomes accessible.**



**By utilizing this existing dataset as a prototype, it provides a roadmap for structuring future data collection endeavors. It allows for the identification of essential variables, relevant features, and potential correlations that might be crucial within the Nigerian context.**

**Moreover, this groundwork helps in establishing methodologies for data preprocessing, outlier detection, feature engineering, and model selection, all of which are fundamental in any agricultural data analysis. These processes can serve as a starting point when integrating future Nigerian agricultural data, ensuring a more streamlined and efficient analysis pipeline.**

**In essence, while the current dataset might not directly represent the Nigerian ecosystem, the methodologies, approaches, and learnings gained through this project create a strong foundation. They stand as a valuable precursor, guiding future data collection efforts and analyses specifically tailored to the unique nuances of the Nigerian agricultural landscape.**

In [ ]: